


Travelling Salesman Problem with Genetic Algorithms

Project Report of Parallel and Distributed Systems: Paradigms
and Models Course
University of Pisa



Antonio Di Mauro
599785

Contents

| | | |
|---|----------------------|---|
| 1 | Introduction | 3 |
| 2 | Workflow analysis | 3 |
| 3 | Speedup analysis | 5 |
| 4 | Experimental results | 6 |
| 5 | Conclusions | 8 |

1 Introduction

Travelling Salesman Problem (TSP) is a problem of finding the shortest path, visiting once the cities/nodes in the graph, from starting city to destination city.

The Genetic Algorithm (GA) is used as heuristic that mimics genetic evolution of species. At the beginning there is a population of chromosomes and at each generation the new population is made by crossover and mutation procedures. Crossover takes two parents and made up a new chromosome/children. Mutation swap randomly two “genes” of a chromosome. The individuals that survives to the next generation is decided by a fitness function that select the best ones.

In the TSP setting a chromosome is a list of cities (identified by an unique integer) that represents a path from a start city to a destination city.

2 Workflow analysis

The problem of TSP with GA could be represented by the pseudocode in Algorithm 1.

Algorithm 1 Pseudocode of TSP with GA

```
create population  
foreach generation:  
    breeding population  
    ranking population
```

At each generation the population size has kepted constant. The breeding operation computes crossover and mutation and the ranking operation sorts the population by fitness.

In Table 1 are collected the times of the single operations obtained running several times the sequential implementation of the algorithm.

| Operation | Latency (μsec) |
|-----------|-----------------------------|
| Creation | 46 |
| Breeding | 86 |
| Ranking | 1628 |

Table 1: Service times of the operations

In a graphical way the workflow of the algorithm, with serial and non serial operations, is represented as in Figure 1.




Figure 1: WorkFlow of the Algorithm (the squared shape represents serial operation and the rounded shape represents the non serial operation)

In the parallel view of the algorithm we could leave out the Creation operation because done just once time at the beginning and focus on the other two.

In a first approach we could think that it's not convenient to parallelize the Breeding operation, because the Ranking operation takes much more time (18 times) than the one said before. So we could think to parallelize only the Ranking operation.

The thought said before is represented using the Master-Worker template of Figure 2.




Figure 2: Master-Worker template with Breeder as Emitter and Ranker as Worker

The Breeder receives the initial population and then computes the breeding operation. After the population is divided in nw (number of workers) sub-populations, these are passed to each Ranker node that computes the ranking operation and returns the sub-population to the Breeder. Finally, the Breeder computes the breeding operation on the top-k (where k is the size of initial population) of the population and discard the others.

The breeding and ranking operations are repeated n times, where n is the number of generations.

The second approach uses RPLSH with optimization/refactoring rules. The

result is that the breeding and the ranking operations are combined in a single node in a configuration with the minimum service time and resources used.

The RPLSH analysis is in Listing 1.

Listing 1: RPLSH analysis

```
rplsh> breed = seq(86)
rplsh> rank = seq(1628)
rplsh> main = pipe(breed,rank)
rplsh> rewrite main with allrules, allrules
rplsh> optimize main with farmopt, pipeopt, maxresources
rplsh> show main by servicetime, resources +5
122.428571      16      [4] : farm(comp(breed,rank)) with [ nw: 14]
122.428571      16      [16] : comp(farm(breed) with [ nw: 14],farm(rank) with [ nw: 14])
125.230769      16      [5] : pipe(breed,farm(rank) with [ nw: 13])
148.000000      16      [1] : pipe(breed,farm(farm(rank) with [ nw: 11]) with [ nw: 1])
148.000000      16      [9] : pipe(farm(breed) with [ nw: 1],farm(rank) with [ nw: 1])
```

The definitive Workflow is shown in Figure 3.

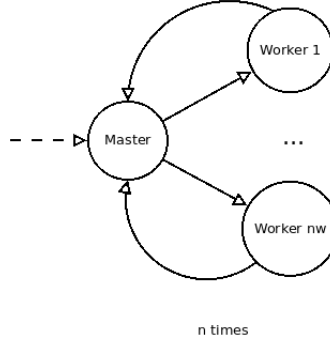


Figure 3: Final Workflow with a Master node that divides and collects the tasks, and Worker nodes that computes breeding and ranking operations

The final workflow is used in the C++ Threads and FastFlow implementations.

3 Speedup analysis

Thanks to Gustafson law we consider the serial fraction that represents the fraction of time spent to executes the Creation operation.

$$f = \frac{t_{serial}}{t_{serial} + t_{non\ serial}} = \frac{t_{creation}}{t_{creation} + n * (t_{breeding} + t_{ranking})} = \frac{46}{46 + n * 1714}$$

the maximum reachable (ideal) speedup is the following one:

$$sp(nw) = \frac{t_{seq}}{f * t_{seq} + [(1 - f) * t_{seq}] / nw}$$

where n is the number of generations and nw is the number of workers.

4 Experimental results

The experiments were performed on two different machines:

1. 2-core with 2-way hyperthreading Dell XPS 13 9360 with Intel(R) Core(TM) i7-7500U CPU @ 2.70GHz \times 4
2. 64-core with 4-way hyperthreading Dell PowerEdge C6320p with Intel(R) Genuine Intel(R) CPU @ 1.30GHz \times 256

| Par. degree | C++ Threads | FastFlow | Ideal |
|-------------|-------------|----------|---------|
| 1 | 0.74727 | 0.880402 | 1 |
| 2 | 0.836531 | 1.74102 | 1.94906 |
| 4 | 0.717611 | 1.49286 | 3.70917 |

Table 2: Speedups on machine #1 (population size=1000, num. of generations=1)

| Par. degree | C++ Threads | FastFlow | Ideal |
|-------------|-------------|----------|---------|
| 1 | 0.752111 | 0.863843 | 1 |
| 2 | 0.893146 | 1.6211 | 1.99466 |
| 4 | 0.725074 | 1.58369 | 3.96814 |

Table 3: Speedups on machine #1 (population size=1000, num. of generations=10)

| Par. degree | C++ Threads | FastFlow | Ideal |
|-------------|-------------|----------|---------|
| 1 | 0.779181 | 0.84016 | 1 |
| 2 | 0.883098 | 1.51184 | 1.99732 |
| 4 | 0.707661 | 1.55068 | 3.98398 |

Table 4: Speedups on machine #1 (population size=1000, num. of generations=20)

| Par. degree | C++ Threads | FastFlow | Ideal |
|-------------|-------------|----------|---------|
| 1 | 0.704687 | 0.837373 | 1 |
| 2 | 0.795344 | 1.63644 | 1.94906 |
| 4 | 0.920407 | 3.24877 | 3.70917 |
| 8 | 1.03912 | 5.96909 | 6.76273 |
| 16 | 1.04937 | 7.1838 | 11.4939 |
| 32 | 0.899429 | 6.14627 | 17.6773 |
| 64 | 0.851438 | 4.15477 | 24.182 |
| 128 | 0.779642 | 2.30888 | 29.6343 |
| 256 | 0.708667 | 1.18282 | 33.3995 |

Table 5: Speedups on machine #2 (population size=1000 and num. of generations = 1

| Par. degree | C++ Threads | FastFlow | Ideal |
|-------------|-------------|----------|---------|
| 1 | 0.690675 | 0.628668 | 1 |
| 2 | 0.825318 | 1.46644 | 1.99466 |
| 4 | 0.933987 | 2.91424 | 3.96814 |
| 8 | 1.05415 | 5.62771 | 7.85287 |
| 16 | 1.16839 | 8.96458 | 15.3824 |
| 32 | 0.862714 | 7.35399 | 29.5482 |
| 64 | 0.821068 | 4.65608 | 54.7652 |
| 128 | 0.775508 | 2.3206 | 95.5275 |
| 256 | 0.767707 | 1.29484 | 152.152 |

Table 6: Speedups on machine #2 with population size = 1000, num. of generations=10)

| Par. degree | C++ Threads | FastFlow | Ideal |
|-------------|-------------|----------|---------|
| 1 | 0.664545 | 0.733771 | 1 |
| 2 | 0.814018 | 1.44799 | 1.99732 |
| 4 | 0.910874 | 2.9926 | 3.98398 |
| 8 | 1.02095 | 5.56504 | 7.92565 |
| 16 | 1.12625 | 8.82946 | 15.6847 |
| 32 | 0.815376 | 7.12631 | 30.7237 |
| 64 | 0.775027 | 4.45778 | 59.0174 |
| 128 | 0.747349 | 2.425 | 109.384 |
| 256 | 0.755664 | 1.23868 | 190.799 |

Table 7: Speedups on machine #2 (population size=1000, num. of generations=20)

5 Conclusions

The C++ Threads implementation faces the problem with the map data parallel pattern, trying to reduce the service time/latency. Its performances are very bad due to the fact that it is not able to amortize the overhead increasing the number of generations.

The FastFlow implementation, instead, faces the problem with the farm stream parallel pattern, trying to increase the throughput. In all the results (on machine #1 and #2) it's obvious that the FastFlow solution performs better than the C++ Threads solution, due to a better speedup. Nevertheless it performs better on machine #2 because of a greater availability of physical cores with respect to machine #1.

Looking the results the problem seems to be difficult to scale because of shorter operations and the overhead does not amortize with increasing number of generations.